

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.datasets
from sklearn.model_selection import train_test_split
from xgboost import XGBRegressor
from sklearn import metrics
```

```
house_price_dataset = sklearn.datasets.fetch_california_housing()
```

```
print(house_price_dataset)
```

```
{'data': array([[ 8.3252, 41., 6.98412698, ..., 2.55555556,
 37.88, -122.23 ],
 [ 8.3014, 21., 6.23813708, ..., 2.10984183,
 37.86, -122.22 ],
 [ 7.2574, 52., 8.28813559, ..., 2.80225989,
 37.85, -122.24 ],
 ...,
 [ 1.7, 17., 5.20554273, ..., 2.3256351,
 39.43, -121.22 ],
 [ 1.8672, 18., 5.32951289, ..., 2.12320917,
 39.43, -121.32 ],
 [ 2.3886, 16., 5.25471698, ..., 2.61698113,
 39.37, -121.24 ]]), 'target': array([4.526, 3.585, 3.521, ..., 0.923, 0.847, 0.894]), 'frame': None, 'target_n:
```

```
# Loading the dataset to a pandas dataframe
```

```
house_price_dataframe = pd.DataFrame(house_price_dataset.data, columns = house_price_dataset.feature_names)
```

```
house_price_dataframe.head()
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25

Next steps:

[Generate code with house_price_dataframe](#)
[View recommended plots](#)
[New interactive sheet](#)

```
# add the target column to the dataframe
```

```
house_price_dataframe['price'] = house_price_dataset.target
```

```
house_price_dataframe.head()
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	price
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422

Next steps:

[Generate code with house_price_dataframe](#)
[View recommended plots](#)
[New interactive sheet](#)

```
# checking the number of rows and columns in the dataframe
```

```
house_price_dataframe.shape
```

```
(20640, 9)
```

```
# check for missing values
```

```
house_price_dataframe.isnull().sum
```



```

pandas.core.frame.DataFrame.sum
def sum(axis: Axis | None=0, skipna: bool=True, numeric_only: bool=False, min_count: int=0,
**kwargs)

>>> pd.Series([np.nan]).sum()
0.0

>>> pd.Series([np.nan]).sum(min_count=1)
nan

```

```

# statistical measures of the dataset
house_price_dataframe.describe()

```



	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	price
count	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	3.870671	28.639486	5.429000	1.096675	1425.476744	3.070655	35.631861	-119.569704	2.068558
std	1.899822	12.585558	2.474173	0.473911	1132.462122	10.386050	2.135952	2.003532	1.153956
min	0.499900	1.000000	0.846154	0.333333	3.000000	0.692308	32.540000	-124.350000	0.149990
25%	2.563400	18.000000	4.440716	1.006079	787.000000	2.429741	33.930000	-121.800000	1.196000
50%	3.534800	29.000000	5.229129	1.048780	1166.000000	2.818116	34.260000	-118.490000	1.797000
75%	4.743250	37.000000	6.052381	1.099526	1725.000000	3.282261	37.710000	-118.010000	2.647250
max	15.000100	52.000000	141.909091	34.066667	35682.000000	1243.333333	41.950000	-114.310000	5.000010

Understanding the correlation between various features in the dataset

Positive Correlation Negative Correlation


```
correlation = house_price_dataframe.corr()
```

```
# constructing a heatmap to understand the correlation
```

```

plt.figure(figsize=(10,10))
sns.heatmap(correlation, cbar=True, square=True, fmt='.1f', annot=True, annot_kws={'size':8}, cmap='Blues')

```



```
X = house_price_dataframe.drop(['price'], axis=1)
Y = house_price_dataframe['price']
```

```
print(X,Y)
```

```

0      MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup  Latitude  \
0      8.3252    41.0    6.984127  1.023810    322.0    2.555556    37.88
1      8.3014    21.0    6.238137  0.971880    2401.0    2.109842    37.86
2      7.2574    52.0    8.288136  1.073446    496.0    2.802260    37.85
3      5.6431    52.0    5.817352  1.073059    558.0    2.547945    37.85
4      3.8462    52.0    6.281853  1.081081    565.0    2.181467    37.85
...      ...      ...      ...      ...      ...      ...      ...
20635   1.5603    25.0    5.045455  1.133333    845.0    2.560606    39.48
20636   2.5568    18.0    6.114035  1.315789    356.0    3.122807    39.49
20637   1.7000    17.0    5.205543  1.120092    1007.0    2.325635    39.43
20638   1.8672    18.0    5.329513  1.171920    741.0    2.123209    39.43
20639   2.3886    16.0    5.254717  1.162264    1387.0    2.616981    39.37

Longitude
0      -122.23
1      -122.22
2      -122.24
3      -122.25
4      -122.25
...      ...
20635   -121.09
20636   -121.21
20637   -121.22
20638   -121.32
20639   -121.24

[20640 rows x 8 columns] 0      4.526
1      3.585
2      3.521
3      3.413
4      3.422
...      ...
20635   0.781
20636   0.771

```

```
20637    0.923
20638    0.847
20639    0.894
Name: price, Length: 20640, dtype: float64
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=2)
```

```
print(X.shape, X_train.shape, X_test.shape)
```

```
(20640, 8) (16512, 8) (4128, 8)
```

```
# load the model
model = XGBRegressor()
```

```
#training the model with X_train
model.fit(X_train, Y_train)
```

```
XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)
```

```
# accuracy for prediction on training data
training_data_prediction = model.predict(X_train)
```

```
print(training_data_prediction)
```

```
[0.5523039 3.0850039 0.5835302 ... 1.9204227 1.952873 0.6768683]
```

```
# R Squared Error
score_1 = metrics.r2_score(Y_train, training_data_prediction)

# Mean Absolute Error
score_2 = metrics.mean_absolute_error(Y_train, training_data_prediction)
```

```
print('R Squared Error:', score_1)
print('Mean Absolute Error:', score_2)
```

```
R Squared Error: 0.943650140819218
Mean Absolute Error: 0.1933648700612105
```

```
plt.scatter(Y_train, training_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Price vs Predicted Price")
plt.show()
```

