

IMAGE CAPTIONING SYSTEM

Presented by: Roopsi Srivastava (202410116100173)

Rani Kumari(202410116100163)

Rabita Yadav(202410116100)

INTRODUCTION

- What is Image Captioning?**

A process of generating textual descriptions for an image using both computer vision and natural language processing.

- Why is it Important?**

- Helps visually impaired people
- Used in content-based image retrieval
- Powers features in social media and e-commerce platforms

Real-World Applications

- Examples:**

- Google Photos automatic captions
- Facebook alt text for visually impaired users
- Pinterest visual search
- Autonomous vehicles understanding their surroundings
- Add 2–3 image examples with captions for effect

How Image Captioning Works

- **High-Level Overview:**
 - **Image Input** → CNN (extract features)
 - **Feature Vector** → RNN/LSTM (generate sentence)
 - **Output** → Caption
- **Visual Aid:** Show a diagram of the above flow.

Components Involved

- **1. CNN (Convolutional Neural Networks):**
Extracts image features (e.g., using VGG, ResNet)
- **2. RNN/LSTM/GRU (Language Model):**
Generates sequence of words based on extracted features
- **3. Attention Mechanism (Optional):**
Focuses on specific parts of the image while generating each word

Algorithms & Models

- **Encoder-Decoder Architecture**
- **Encoder:** CNN
- **Decoder:** RNN or LSTM with Word Embeddings
- **Pretrained Models:**
 - InceptionV3, ResNet50 (for image encoding)
 - Beam Search or Greedy Decoding (for caption generation)

Tools and Technology

- **Languages:** Python
- **Libraries:** TensorFlow / PyTorch, OpenCV, NLTK
- **Frameworks:** Keras, HuggingFace Transformers
- **Other:** Google Colab

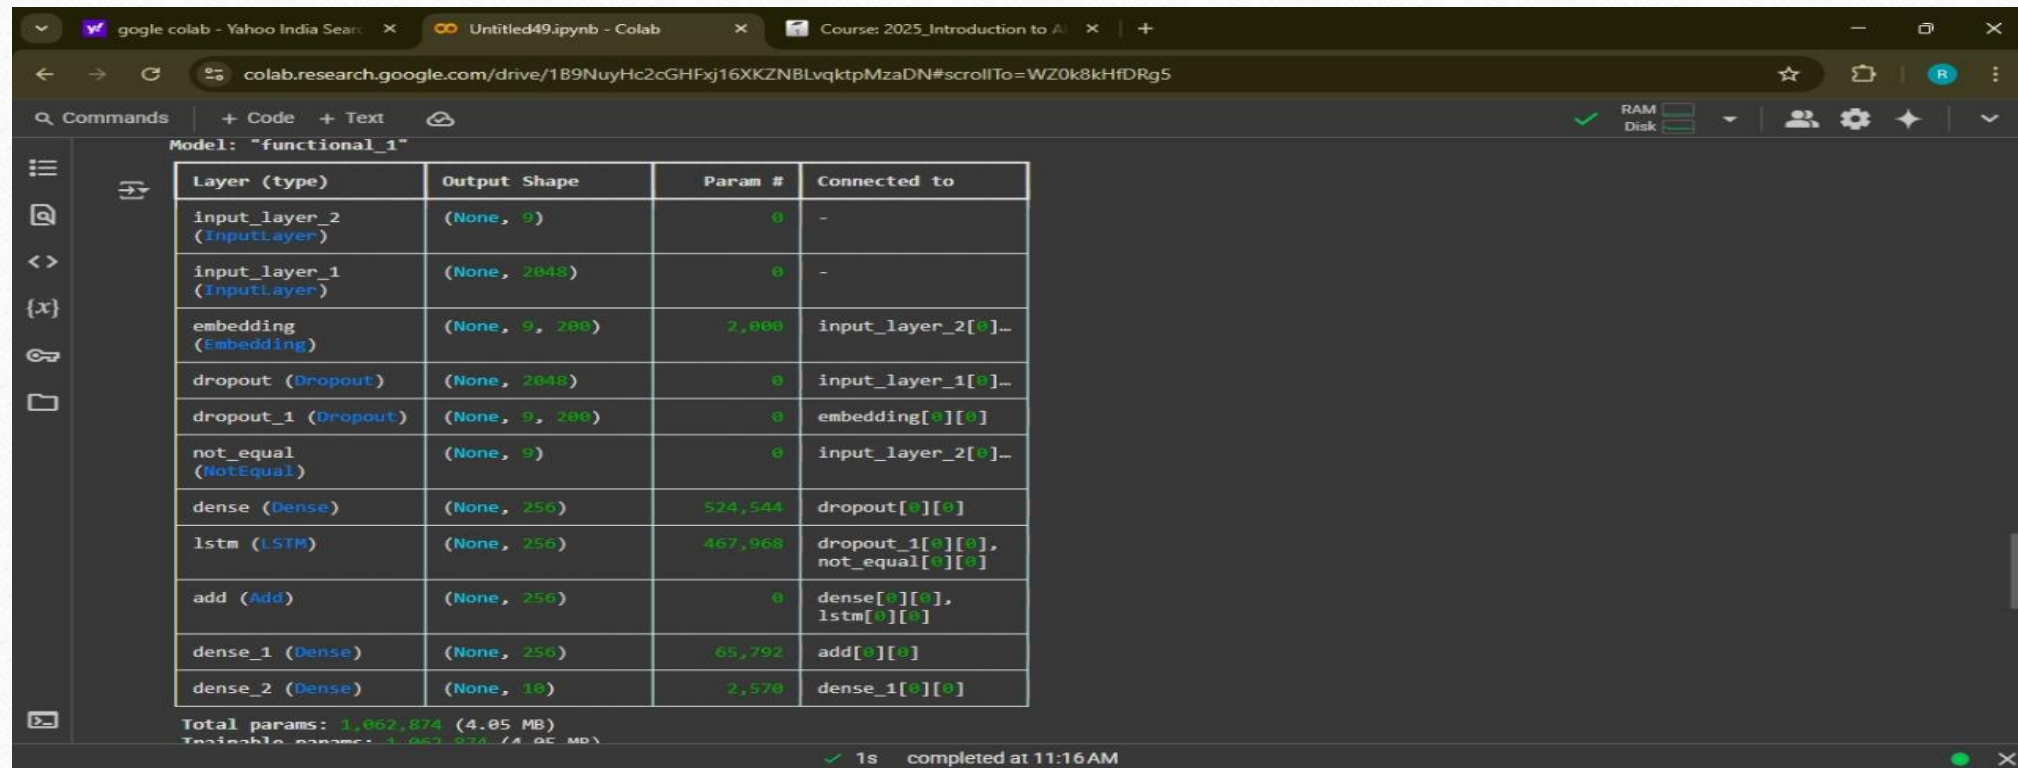
Future Scope

- Improved caption quality with GPT models
- Multilingual caption generation
- Real-time captioning in AR/VR
- Personalized image descriptions

Image Uploaded



Output



The screenshot displays a Google Colab environment with a browser window showing the Colab interface. The main content area displays the summary for a model named "functional_1". The summary is presented as a table with columns: Layer (type), Output Shape, Param #, and Connected to. The layers listed are: input_layer_2 (InputLayer), input_layer_1 (InputLayer), embedding (Embedding), dropout (Dropout), dropout_1 (Dropout), not_equal (NotEqual), dense (Dense), lstm (LSTM), add (Add), dense_1 (Dense), and dense_2 (Dense). The total number of parameters is 1,062,874 (4.05 MB). The training process is shown as completed at 11:16 AM.

Layer (type)	Output Shape	Param #	Connected to
input_layer_2 (InputLayer)	(None, 9)	0	-
input_layer_1 (InputLayer)	(None, 2048)	0	-
embedding (Embedding)	(None, 9, 200)	2,000	input_layer_2[0]...
dropout (Dropout)	(None, 2048)	0	input_layer_1[0]...
dropout_1 (Dropout)	(None, 9, 200)	0	embedding[0][0]
not_equal (NotEqual)	(None, 9)	0	input_layer_2[0]...
dense (Dense)	(None, 256)	524,544	dropout[0][0]
lstm (LSTM)	(None, 256)	467,968	dropout_1[0][0], not_equal[0][0]
add (Add)	(None, 256)	0	dense[0][0], lstm[0][0]
dense_1 (Dense)	(None, 256)	65,792	add[0][0]
dense_2 (Dense)	(None, 10)	2,570	dense_1[0][0]

Total params: 1,062,874 (4.05 MB)
Trainable params: 1,062,874 (4.05 MB)

✓ 1s completed at 11:16 AM

