

HOUSE PRICE PREDICTION

**A PROJECT REPORT
for
Introduction to AI (AI101B)
Session (2024-25)**

Submitted by

**Tanishka Gupta (202410116100218)
Shipra Upadhyay (202410116100196)
Vanshika Garg (202410116100235)
Tripti Rajpoot (202410116100225)**

**Submitted in partial fulfilment of the
Requirements for the Degree of**

MASTER OF COMPUTER APPLICATION

**Under the Supervision of
Mr. Apoorv Jain
Assistant Professor**



Submitted to

**DEPARTMENT OF COMPUTER APPLICATIONS
KIET Group of Institutions, Ghaziabad
Uttar Pradesh-201206
(MAY 2025)**

CERTIFICATE

Certified that **Tanishka Gupta 202410116100218, Shipra Upadhyay 202410116100196, Vanshika Garg 202410116100235 and Tripti Rajput 202410116100225** have carried out the project work having “**HOUSE PRICE PREDICTION**” (INTRODUCTION TO AI (AI101B)) for **Master of Computer Application** from Dr. A.P.J. Abdul Kalam Technical University (AKTU) (formerly UPTU), Lucknow under my supervision. The project report embodies original work, and studies are carried out by the student himself/herself and the contents of the project report do not form the basis for the award of any other degree to the candidate or to anybody else from this or any other University/Institution.

Mr. Apoorv Jain

Assistant Professor

Department of Computer Applications

KIET Group of Institutions, Ghaziabad

Dr. Akash Rajak

Dean

Department of Computer Applications

KIET Group of Institutions, Ghaziabad

HOUSE PRICE PREDICTION

Tanishka Gupta, Shipra Upadhyay, Vanshika Garg, Tripti Rajput

ABSTRACT

A Machine Learning-based approach was adopted for this project to help predict house prices through regression techniques and feature engineering. Accurate price prediction in the real estate market helps buyers, sellers, and investors to take decisions based on data. The dataset contains several attributes of the residential property, both numerical (area, number of rooms, construction year, etc.) and categorical (neighborhood, building type, quality, etc.).

Luxuriating in thorough data preprocessing such as dealing with missing values through imputation techniques, encoding categorical variables using One-Hot Encoding, and scaling features by standardization quite helped boost the performance of the models. During all this, the dataset was made clean and consistent to facilitate the requirements of machine learning algorithms.

Three regression models were implemented and compared as part of the project: **Linear Regression, Ridge Regression, and Lasso Regression**. After splitting the datasets in various ways, the models were trained, tested, and their results evaluated in terms of RMSE. Because of the regularization, the Ridge and Lasso were with very less chances of overfitting and would produce better generalizations.

Visualization and residual analyses were carried out to understand model behaviors and the accuracy of predictions. Based on the experiment, regularized regression methods scored well in stability, and accuracy.

The project underscores the thrust given to machine learning and data preprocessing in real- life problems. It states that when proper feature engineering and model building approach are considered; it is possible to build predictive systems that can fairly convince an estimation of a property's price.

ACKNOWLEDGEMENTS

Success in life is never attained single-handedly. My deepest gratitude goes to my project supervisor, **Mr. Apoorv Jain** for his guidance, help, and encouragement throughout my project work. Their enlightening ideas, comments, and suggestions.

Words are not enough to express my gratitude to **Dr. Akash Rajak, Professor and Dean**, Department of Computer Applications, for his insightful comments and administrative help on various occasions.

Fortunately, I have many understanding friends, who have helped me a lot on many critical conditions.

Finally, my sincere thanks go to my family members and all those who have directly and indirectly provided me with moral support and other kind of help. Without their support, completion of this work would not have been possible in time. They keep my life filled with enjoyment and happiness.

Tanishka Gupta

Shipra Upadhyay

Vanshika Garg

Tripti Rajput

TABLE OF CONTENTS

Certificate	2
Abstract	3
Acknowledgements	4
Table of Contents	5
1. Introduction	7-11
1.1 Background on House Prediction	7
1.2 Scope of the Project	8
1.3 Purpose and Objectives	9
1.4 Significance	10
1.5 Significance	11
2. Feasibility Study	12-13
2.1 Technical Feasibility	12
2.2 Economical Feasibility	12
2.3 Operational Feasibility	13
2.4 Legal Feasibility	13
2.5 Timeline Feasibility	13
3. Project Objectives	14-17
3.1 Key Objectives	14-15
3.2 Project Objectives	16-17
4. Technology	18-19
5. Project Flow	20-28

5.1 Problem Description	20
5.2 Requirement Analysis	21
5.3 Design Phase	22-23
5.4 Technology Selection	24-25
5.5 Development	26-27
5.6 Testing and Validation	28
6. Project Outcome	29-34
6.1 Correlation Heatmap	30
6.2 Distribution of Sale Prices	31
6.3 Actual vs Predicted Prices	32
6.4 Neighborhood-wise Sale Price Distribution	33
6.5 Residual Distribution	34
7. References	

CHAPTER 1

INTRODUCTION

1.1 Background on House Prediction

The real estate sector is an economic engine and one of the largest contributors to the economy with innumerable stakeholders like buyers, sellers, investors, and financial institutions. Among the most crucial factors in the real estate decision-making is ascertaining the fair market value of a given property. Traditionally, valuation of a property was performed by real estate agents or appraisers using their personal experience and knowledge of location as well as comparable sales. Yet, in all fairness, this process has always been thinking with the stake of subjectivity, given time constraints, and prone to human error. However, with the evolution of AI and data-driven technologies, the world of property valuation is witnessing a metamorphosis.

House Price Prediction (HPP) uses machine-learning algorithms and statistical techniques that forecast future price by analysing past sales data. It factors in many attributes such as location, size, number of rooms, facilities, age of the building, and neighbourhood demographics toward obtaining accurate price estimations. The primary objective is to reduce any uncertainty or subjectivity that comes with manual valuation by trying computational models that actually learn from the data.

This technological development has caused house price prediction to be swifter, more reliable, and accessible. Hence, this helps individuals who want to make decisions either to buy or sell their house, or banks that want to assess the risk of mortgages, or developers that want to look at the trend of this market. Given the enormous accumulation of housing data and AI's capability to decipher complicated patterns, this lays the groundwork for an important application of predictive analytics. With this, advance price prediction of houses by AI is now an integral component of modern real estate and financial decision-making processes.

1.2 Scope of the Project

General scope of the House Price Prediction project involves building a strong and scalable machine learning model that can estimate residential property selling prices based on many input features. The emphasis of this project lies within employing structured historical data that consists of numerical as well as categorical variables to train the models so that they behave well and generalize on new housing data.

Within the scope fall areas such as data collection, data preprocessing (treating missing values, encoding categorical features, scaling), feature engineering, model training, and evaluation. This project uses models such as **Linear Regression**, **Ridge Regression**, and **Lasso Regression**, which consider continuous value prediction as an application. Employing the regularized regression models helps with multicollinearity or overfitting.

The dataset usually consists of attributes like size of the property, bedrooms, bathrooms, location, construction year, overall quality, and features such as garage space or distance to amenities. This project is limited to structured tabular data and does not include unstructured inputs like images or textual property descriptions.

Whereas the existing scope is concerned with developing a proof-of-concept system that operates with historical datasets, it may further be extended by incorporating real-time data streams, geographic information systems (GIS), or deep models such as Gradient Boosting and Neural Networks for better accuracy.

The project is intended to be done with Python and open-source packages such as Pandas, Scikit-learn, Matplotlib, and Seaborn and hence is accessible, reproducible, and simple to deploy for academic or applied purposes of real estate valuation.

1.3 Purpose and Objectives

The primary purpose of the House Price Prediction project is to develop an intelligent, data-driven system that can accurately predict the selling prices of residential properties based on their features. The system aims to assist stakeholders in the real estate ecosystem—including buyers, sellers, investors, agents, and financial institutions—by providing reliable and consistent property valuations that reduce the reliance on subjective judgments. The overarching goal is to harness the power of machine learning to uncover patterns and relationships within housing data that might not be immediately apparent through traditional analysis. By doing so, the project seeks to create a predictive model that is both accurate and generalizable to real-world housing markets. To achieve this purpose, the following objectives are defined:

- **Data Preparation and Understanding:** Obtain and investigate a housing dataset to see what its structure is, have missing or inconsistent values, and make it ready for analysis with preprocessing and cleaning.
- **Feature Engineering:** Choose and reformat the most important features that affect house prices, such as location, size, facilities, and property age.
- **Model Development:** Train and implement several regression models, including Linear Regression, Ridge, and Lasso, to find the best method for price prediction.
- **Model Evaluation:** Utilize suitable measures such as Root Mean Squared Error (RMSE) to measure model precision and test its performance on unseen test data.
- **Insights and Interpretation:** Present qualitative insights from the model outputs and interpret the role of important features in the prediction.

1.4 Significance

House Price Prediction is a valuable project both academically and practically. In a world where there is plenty of data and decisions have to be made quickly and accurately, using machine learning for estimation of real estate prices has many benefits. Conventional ways of property valuation are prone to human prejudice, disparity, and area-based subjectivity. This project tries to minimize such constraints by presenting a data-based, homogeneous, and scalable method of property pricing.

Academically, the project presents an **end-to-end** application of machine learning principles—spanning from **data preprocessing** and **feature engineering** to model assessment and optimization. It is a practical exercise exploring how regression methods can be utilized to address real-world issues and is thus a great learning resource for students and researchers in AI, data science, and real estate analysis.

In practice, the model can be used by multiple stakeholders. Buyers are able to have an objective idea of **fair market value**, sellers can list their homes more competitively, and investors are able to determine undervalued or overvalued properties. Financial institutions and mortgage lenders are also able to use such models in order to measure risk prior to lending approval.

In addition, AI integration in real estate can help **urban planning** and **policy-making** by analysing market trends and regional price variances. Since the model is capable of making instant predictions, it is ideal for real-time applications like property listing websites or mobile apps.

In general, the importance of the project is that it holds the capability to transform property valuation with its **accuracy, accessibility, and efficiency**—closing the loopholes between conventional appraisal techniques and new technology.

1.5 Overview of Methodology

The approach utilized in the House Price Prediction project is designed to conform to the standard machine learning pipeline to ensure simplicity, scalability, and precision throughout the predictive process. The project starts with **data acquisition**, employing an open-source housing dataset with multiple features like number of rooms, square feet, location, house age, and other pertinent factors affecting house prices.

The second step is **data preprocessing** wherein missing values are addressed through proper imputation methods, categorical variables are converted through One-Hot Encoding, and numerical attributes are **standardized** to have uniformity. This step makes the data clean and properly formatted for machine learning algorithms.

After preprocessing, **exploratory data analysis (EDA)** is carried out to know data distribution, detect correlations, and view important patterns. EDA facilitates **feature selection and engineering**, in which only the most relevant attributes are kept or restructured to improve model performance.

Later, the data is divided into training and test subsets, and several **regression models** such as **Linear Regression**, Ridge Regression, and Lasso Regression are constructed. They are trained on the training data and tested on the test data with Root Mean Squared Error (RMSE) as the key performance metric. Regularized models prevent overfitting and enhance generalization.

Lastly, **model evaluation** and **residual analysis** are executed to measure prediction accuracy and error behaviour. The top-performing model is used for deployment. This systematic approach ensures the prediction system is both efficient and realistic for real estate applications in practice.

CHAPTER 2

FEASIBILITY STUDY

A feasibility study analyses the practicality and viability of putting a proposed system in place by analysing various critical dimensions. For the House Price Prediction system, five broad areas—technical, economic, operational, legal, and timeline feasibility—are analysed to identify if the project can be developed and rolled out successfully.

2.1 Technical Feasibility

The technical feasibility assesses whether or not the requisite technology, tools, and skills exist to create the project. This framework utilizes open-source programming languages such as **Python** and libraries including **Pandas**, **NumPy**, **Scikit-learn**, **Matplotlib**, and **Seaborn** for preprocessing, visualization, and machine learning. These are well-supported and are available free of cost. The models used—**Linear Regression**, **Ridge Regression**, and **Lasso Regression**—are efficient from a computational perspective and can be executed on a normal computer with reasonable specs (i.e., 8GB RAM, i5 processor). Hence, no specialized infrastructure is required. With a good grasp of data science fundamentals and availability of the necessary software and dataset, the project is technically viable.

2.2 Economical Feasibility

The economic feasibility takes into account the cost of the project. As this project is more theoretical and based on open-source tools and public datasets (e.g., Kaggle house prices), the cost of development is low. No costly licenses, cloud resources, or third-party APIs are required. The primary investment is effort and time, with zero direct money costs. In future development as a commercial product, hosting, API services, or high-end datasets could be costs. None of these would be possible without plenty of time and effort, though. At present, though, the project can be costed for an economic student or university.

2.3 Operational Feasibility

Operational feasibility looks at whether the system will work in the target environment and also be used by the users. The suggested system for predicting is easy to use, interpretable, and produces intelligible results. With a minimal user interface or integration into existing platforms (e.g., a real estate portal), end users can enter house information and obtain price estimates. This assists stakeholders in real estate transactions and market analysis. Since it doesn't involve specialized know-how for utilizing, the system is operationally viable and could be implemented broadly with minimal user training or guidelines.

2.4 Legal Feasibility

Legal feasibility addresses any data usage or AI implementation-related regulatory or compliance issues. The project utilizes *open-source housing datasets* and adheres to data usage rules, preventing it from infringing any copyright or privacy act. There is no sensitive personal information involved. If the model were to be used commercially, legal compliance with legislation such as GDPR or consumer data law may become a requirement—particularly in case of integration of user-generated or location-based data. But in the case of academic usage, there are no legal restrictions involved, so this feature is possible.

2.5 Timeline Feasibility

Timeline feasibility determines whether it is possible to finish the project within given time. With well-defined stages - data preprocessing, model training, evaluation, and documentation—the project can be accomplished within 6 to 8 weeks with continuous progress. Every stage takes 1 to 2 weeks based on complexity. As most tools already exist and there are no external dependencies, the latency is negligible. With good planning, the timeline is viable and practical, which means that the project can be accomplished within a typical academic term or project cycle.

CHAPTER 3

PROJECT OBJECTIVE

3.1 Key Objectives

The overall objectives of the House Price Prediction project centre on creating an interpretable, accurate, and efficient machine learning system capable of predicting residential property prices based on input factors. The main objectives are mentioned below:

1. Develop a Predictive Model:

Create a machine learning model with regression algorithms (Linear, Ridge, and Lasso) to forecast house prices given a set of property features.

2. Ensure High Model Accuracy:

Reduce prediction error with performance measures like RMSE (Root Mean Squared Error) and enhance model generalization with regularization methods.

3. Efficient Data Preprocessing:

Preprocess the dataset by resolving missing values, encoding categorical variables, and feature scaling to make the data uniform and model compatible.

4. Execute Exploratory Data Analysis (EDA):

Examine and plot data trends, distributions, and correlations between features for deeper understanding prior to model training.

5. Feature Selection and Engineering

Identify the most significant factors impacting house prices and create new variables if needed to enhance predictive capability.

6. Compare Multiple Models:

Train and compare various regression models to find out which model works best under the current dataset.

7. Interpret Model Outputs:

Offer interpretability by describing how varying features contribute to the end price prediction, making the system comprehensible and reliable.

8. Design for Reusability and Scalability:

Ensure the model and code are organized in a modular way, so updates or improvements in the future can be added with ease.

9. Create a Prototype Interface (Optional):

If time allows, create a simple user interface or command-line tool for end-users to enter house information and see forecasted prices.

10. Document the Process:

Keep accurate records of the process, code, and results for educational, training, or development use.

3.2 Project Objectives

The House Price Prediction project has specific objectives that intend to solve particular problems in real estate valuation through artificial intelligence and machine learning. The subsequent comprehensive objectives reflect what the project will accomplish during its development cycle:

- 1. Accurate Price Estimation:**

Develop a machine learning model capable of accurately forecasting house prices by learning from past data.

- 2. Utilization of Regression Techniques:**

Run and compare several regression algorithms like Linear Regression, Ridge Regression, and Lasso Regression to determine the top-performing method.

- 3. Data-Driven Insights:**

Apply data analytics to find the major factors affecting housing prices, viz., location, size of the property, amenities, and quality of construction.

- 4. Model Optimization:**

Enhance model performance using hyperparameter tuning, regularization, and feature selection to prevent overfitting and achieve generalization on unobserved data.

- 5. Handling Complex Datasets:**

Clean missing values, encode categorical variables, and normalize numerical features to prepare complex real-world datasets.

- 6. Simplifying Decision-Making:**

Offer a system that can help stakeholders—buyers, sellers, and real estate agents—make informed decisions on the basis of dependable price estimates.

7. Model Interpretability:

Make the model interpretable by emphasizing how each feature affects the final prediction, which fosters transparency and trust.

8. Prototype Development:

Optionally create an easy-to-use interface that enables users to enter property characteristics and receive price estimates in real-time.

9. Validation and Testing:

Test the model using criteria like **Root Mean Squared Error (RMSE)**, **R² Score**, and cross-validation in order to provide reliability and stability.

10. Knowledge Enhancement:

Enhance understanding of machine learning applications in real-world scenarios and contribute valuable insights for academic or practical use in the domain of real estate.

CHAPTER 4

TECHNOLOGY

House Price Prediction project is developed with the latest, efficient, and open-source technologies. The technologies have been selected with care to ease development, make precise predictions, and facilitate collaboration.

1. Python

Python is the main programming language employed in this project. It is widely used in data science and machine learning because of its straightforward syntax, readability, and rich set of effective libraries. Python facilitates easy handling of big data, implementation of algorithms, and visualization of results effectively.

2. Libraries and Frameworks

In order to develop a full-fledged machine learning pipeline, the following Python libraries and frameworks were utilized:

- **Pandas:** For data cleaning and manipulation. It is useful for managing missing values, joining datasets, and feature creation.
- **NumPy:** Offers numerical array and operation support. It plays a fundamental role in efficient handling of high-dimensional data.
- **Matplotlib and Seaborn:** These two libraries were employed for data visualization. They assist in creating charts, graphs, and heatmaps to identify patterns and relationships within the data.
- **Scikit-learn (sklearn):** A fundamental machine learning library utilized for the implementation of regression algorithms such as Linear Regression, Ridge Regression, and Lasso Regression. It further includes functionalities to evaluate a model, preprocess data, and validate.

3. Google Colab

Google Colab is a web-based interface where a user can write and execute Python code in Jupyter Notebook format. It is particularly helpful for machine learning projects due to the following:

- No installation or setup is needed.
- It provides free access to strong computing capabilities such as GPUs.
- It supports real-time collaboration, with several users being able to work on a single notebook.
- is integrated directly with Google Drive, providing seamless file sharing and access.

CHAPTER 5

PROJECT FLOW

5.1 Problem Description

Nowadays, it is a crucial but difficult task to accurately forecast house prices in the property market. Using conventional methods and tools is not satisfactory because of many inherent shortcomings:

1. **Complicated Market Factors:** House prices are influenced by numerous factors such as location, size of the property, age, local facilities, economic environment, and prevailing market trends, making it hard to estimate the price.
2. **Data Overload:** The increasing amount and diversity of real estate data may exceed the capability of manual or straightforward statistical analysis methods, resulting in inaccurate findings.
3. **Rigid Models:** Most current price forecasting models are based on linear or naive assumptions that are unable to detect nonlinear relationships or intricate feature interactions.
4. **Lack of Access to Advanced Tools:** Advanced predictive tools are expensive and out of reach for average buyers, sellers, or small real estate agents.
5. **No Dynamic Real-Time Updates:** Housing markets keep changing, but most systems do not offer dynamic, real-time price forecasts based on current patterns.
6. **Time-Consuming and Error-Prone Processes:** Pricing analysis is time-consuming and prone to errors without automation, creating decision delays.

This project solves these challenges by creating an AI-powered house price estimation system that incorporates machine learning algorithms to scan massive datasets, identify sophisticated patterns, and provide precise, customized, and timely price estimates that are accessible to everyone involved.

5.2 Requirements Analysis

In order to create a successful AI-based house price predictor, the following are the most important requirements that should be addressed:

1. **Data Gathering:** Collect complete and good quality real estate information such as property attributes (size, location, number of rooms), economic factors, previous prices, neighbourhood information, and market trends.
2. **Data Preprocessing:** Manage missing values, eliminate outliers, normalize and encode the categorical variables to make the data suitable for machine learning algorithms.
3. **Feature Selection:** Determine the key features most significantly affecting house prices to enhance model efficiency and lower computational complexity.
4. **Model Development:** Adopt machine learning models that can infer intricate, nonlinear relationships, like Random Forest, Gradient Boosting, or Neural Networks.
5. **User Interface:** Design an easy-to-use interface through which users can provide property information and obtain estimated price predictions in a timely manner.
6. **Real-time Prediction:** Enable the system to update predictions dynamically based on recent market data or user inputs.
7. **Scalability:** Design the system to handle large datasets efficiently and support multiple concurrent users without performance degradation.
8. **Accuracy and Validation:** Establish rigorous testing and validation protocols to ensure prediction accuracy, using techniques like cross-validation and performance metrics (RMSE, MAE).
9. **Security and Privacy:** Protect sensitive user data and ensure compliance with data privacy regulations.
10. **Accessibility:** Make the system accessible by different user groups for house buyers, sellers, agents, or investors by web and mobile platforms.

5.3 Design Phase

The design phase outlines the plan for developing the AI house price prediction system, with emphasis on architecture, data flow, and user interface:

1. System Architecture:

The system follows a modular architecture with data ingestion, preprocessing, model training, prediction engine, and user interface layers to ensure scalability and maintainability.

2. Data Flow Design:

Raw real estate data is first collected from multiple sources (public datasets, real estate portals, APIs). It then flows through a preprocessing pipeline for cleaning, transformation, and feature engineering before feeding into the machine learning model.

3. Machine Learning Model Design:

The architecture involves choosing suitable algorithms (e.g., Random Forest, XGBoost) and hyperparameter tuning techniques. The model will support regression tasks and output confidence intervals along with predictions.

4. User Interface Design:

The UI will be designed to be responsive and intuitive in order to enable users to simply enter property information such as location, size, and facilities. It will show predicted prices along with reasons for transparency.

5. Database Design:

A strong database schema is established to hold historical property data, user input, prediction output, and logs. Optimized indexing and querying are given the top priority to achieve quick response times.

6. API Design:

RESTful APIs are created to provide the interface for communication between the frontend and backend to allow real-time data sharing and prediction requests.

7. Security Design:

User authentication, authorization, and data encryption mechanisms are incorporated to protect sensitive information and ensure privacy compliance.

8. Scalability and Performance:

The design considers load balancing, caching, and asynchronous processing to support multiple users and large datasets without compromising speed.

5.4 Technology Selection

Technology selection is important for ensuring the scalability, efficiency, and performance of the house price prediction system. Following is a breakdown of the selected technologies for different components of the project:

1. Programming Language:

- Python is chosen because of its high support for machine learning, data analysis, and scientific computing.
- Data handling and visualization libraries like **NumPy**, **Pandas**, and **Matplotlib** are used.

2. Machine Learning Frameworks:

- Scikit-learn for general regression model implementation like Linear Regression, Random Forest, and Decision Trees.
- XGBoost and LightGBM are utilized for gradient boosting because they offer high accuracy and performance in prediction problems.
- TensorFlow or Keras can be used in case deep learning models (such as neural networks) are investigated.

3. Data Storage:

- **SQLite** or **PostgreSQL** for storing structured data.
- CSV/Excel files at initial development and test phases for ease.
- Cloud-based storage (such as AWS S3, Google Drive) can be used for large-scale deployments.
- Web Development: Flask or Django as backend frameworks for managing API requests and model integration.
- **HTML**, **CSS**, **JavaScript** for frontend.
- React.js can be utilized for developing an interactive and dynamic user interface.

4. APIs & Integration:

- RESTful APIs are built for seamless communication between frontend and backend.
- APIs can also be employed to retrieve real-time housing data if available.

5. Deployment & Hosting:

- **Heroku, Render, or AWS EC2** to host the application.
- Docker for containerization to create uniform environments at deployment.

6. Version Control:

- **Git** and **GitHub** for source code control, team collaboration, and versioning.

5.5 Development

The development stage involves bringing the design and requirements into a working, AI-based house price prediction system. The stage is broken down into various phases to ensure that the build process is smooth and effective:

1. Environment Setup:

- Installed required tools and libraries such as Python, Jupyter Notebook, Scikit-learn, Pandas, NumPy, and Flask. Set up virtual environments to keep dependencies clean.

2. Data Acquisition and Preprocessing:

- Gathered housing data sets from reputable sources such as Kaggle or government real estate websites.
- Conducted data cleaning: missing values, duplicates removal, incorrect data type corrections.
- Created new features (e.g., price per sqft, property age).
- Encoded categorical variables via Label Encoding and One-Hot Encoding.

3. Model Building:

- Trained various regression models like Linear Regression, Random Forest, and XGBoost.
- Optimized hyperparameters by Grid Search and Cross Validation for improved performance.
- Chose the model with the best performance out of RMSE and R^2 score.

4. Backend Development:

- Developed a Flask-based API to host the machine learning model and get input data from the frontend.
- Used endpoints to accept user data and deliver predictions.

5. Frontend Development:

- Built a minimalist and clean UI through HTML, CSS, and optionally React.js.
- Implemented input fields for size, location, rooms, etc., and dynamically showed predicted house prices.

6. Integration:

- Integrated frontend with backend APIs to provide end-to-end functionality.
- Checked outputs against user query samples.

5.6 Testing and Validation

Validation ensures that the machine learning model generalizes well on unseen data and maintains accuracy across various input scenarios. Below are the key validation strategies employed in this project:

1. Data Splitting for Validation:

The dataset was split into training (80%) and testing (20%) subsets.

This method ensures the model is trained on one portion and validated on a separate, unseen portion for unbiased performance metrics.

2. K-Fold Cross Validation:

Applied **K-Fold Cross Validation** (commonly with $k = 5$ or 10) to assess model performance more reliably.

The data is split into k parts; the model is trained on $(k-1)$ parts and validated on the remaining part, repeating this k times.

This technique reduces variance and ensures that every data point is used for both training and validation.

3. Evaluation Metrics for Validation:

To assess the model's performance during validation, the following metrics were used:

- **Mean Absolute Error (MAE):** Measures average absolute prediction error.
- **Root Mean Squared Error (RMSE):** Penalizes larger errors more than MAE.
- **R^2 Score (Coefficient of Determination):** Indicates how well the predictions approximate the actual data.

4. Hyperparameter Tuning:

Grid Search CV was used in conjunction with cross-validation to find the optimal hyperparameters for models like Random Forest and XGBoost.

5. Result Interpretation:

Validation results guided the selection of the final model. XGBoost showed the best generalization capability with high R^2 and low RMSE.

CHAPTER 6

PROJECT OUTCOME

The House Price Prediction project successfully achieved its primary objectives by developing an AI-based system capable of estimating property prices with high accuracy and efficiency. The key outcomes are outlined below:

1. Accurate Predictions:

- The system accurately predicts house prices based on input features like location, size, number of rooms, and other relevant parameters.
- The best-performing model (e.g., XGBoost or Random Forest) demonstrated low RMSE and high R^2 scores, validating its reliability.

2. Smart Model Integration:

- Machine learning models were seamlessly integrated into a user-friendly web application, allowing real-time interaction and instant results for users.

3. User-Friendly Interface:

- The application features an intuitive interface for both technical and non-technical users, making it easy to input data and interpret results.
- Aesthetic design and responsive layout enhance user experience.

4. Scalable and Modular Design:

- The system architecture supports easy upgrades and scalability, such as incorporating more data, new features, or additional machine learning models.


5. Real-World Usability:

- Users such as buyers, sellers, and real estate agents can leverage the platform to make informed decisions backed by data-driven insights.
- The solution bridges the gap between traditional property valuation methods and modern AI-based analytics.

6. Learning and Skill Development:


- The project helped strengthen skills in machine learning, data preprocessing, model evaluation, web development, and system integration.

6.1 Correlation Heatmap




- Variables most positively correlated with SalePrice.
- Top features: OverallQual (0.81), GrLivArea (0.70), GarageCars, GarageArea, and TotalBsmtSF.
- Quality and living area have the strongest influence on house prices.
- Garage-related features also have strong correlations — may indicate luxury or space importance.

6.2 Distribution of Sale Prices




- Right-skewed distribution.
- Most houses are priced between \$100,000 and \$250,000.
- Long tail shows luxury/high-end outliers.
- Log-transforming SalePrice may improve model performance due to skewness.
- Many typical homes cluster around the lower-middle price range.

6.3 Actual vs Predicted Prices (Scatter Plot)




- A near-linear relationship between actual and predicted values.
- Model performs well, especially in the \$100k–\$400k range.
- Some deviation in high-priced houses.
- The model generalizes well, though it struggles with extreme values.
- More high-end training data or outlier handling might help.

6.4 Neighborhood-wise Sale Price Distribution (Box Plot)



- Significant price variation across neighborhoods.
- NoRidge, NridgHt, and StoneBr have higher medians and upper ranges.
- Neighborhood is a critical feature — likely one-hot encoded in your model.
- Price trends are geographically dependent; may indicate school zone, amenities, etc.

6.5 Residual Distribution



- Bell-shaped distribution centered around 0.
- Slight skew indicates mild over- or under-prediction in some cases.
- Model errors are approximately normally distributed — good sign.
- Suggests the model is unbiased overall with consistent performance.

REFERENCES

- Géron, A. (2019). Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow (2nd ed.). O'Reilly Media.
- Raschka, S., & Mirjalili, V. (2019). Python Machine Learning (3rd ed.). Packt Publishing.
- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). An Introduction to Statistical Learning: with Applications in R (2nd ed.). Springer.
- Scikit-learn Developers. (2023). Scikit-learn: Machine Learning in Python. Retrieved from: (<https://scikit-learn.org/>)
- Chen, T., & Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. In Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (pp. 785–794).
- Kaggle. (n.d.). House Prices: Advanced Regression Techniques Dataset. Retrieved from: (<https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques>)
- Brownlee, J. (2020). How to Evaluate Machine Learning Models. Machine Learning Mastery. Retrieved from: (<https://machinelearningmastery.com/how-to-evaluate-machine-learning-models/>)
- TensorFlow Developers. (2023). TensorFlow: An end-to-end open-source machine learning platform. Retrieved from: (<https://www.tensorflow.org/>)
- Pedregosa, F., Varoquaux, G., Gramfort, A., et al. (2011). Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research, 12, 2825–2830.

House Price Prediction

1st Tanishka Gupta
Master of Computer Applications
KIET Group of Institutions
Ghaziabad, India
tanishka.2426mca714@kiet.edu

2nd Shipra Upadhyay
Master of Computer Applications
KIET Group of Institutions
Ghaziabad, India
shipra.2426mca528@kiet.edu

3rd Vanshika Garg
Master of Computer Applications
KIET Group of Institutions
Ghaziabad, India
vanshika.2426mca715@kiet.edu

4th Tripti Rajput
Master of Computer Applications
KIET Group of Institutions
Ghaziabad, India
tripti.2426mca319@kiet.edu

5th Mr. Apoorv Jain
Master of Computer Applications
KIET Group of Institutions
Ghaziabad, India
apoorv.jain@kiet.edu

Abstract—This project aims to predict house prices using machine learning algorithms based on factors like location, size, number of rooms, and property age. Various models such as Linear Regression, Decision Tree, and Random Forest are implemented and evaluated using metrics like RMSE and R-squared. After preprocessing and analyzing a real-world dataset, Random Forest achieved the highest accuracy. The study demonstrates the effectiveness of machine learning in providing reliable and automated solutions for real estate price estimation. The real estate market is influenced by various factors such as location, property size, and local amenities. Traditional methods of price estimation often lack accuracy and adaptability. With advancements in machine learning, predictive models now offer reliable solutions for estimating house prices. This project aims to develop a system that uses historical housing data to forecast property prices accurately. By analyzing key features using algorithms like Linear Regression and Decision Trees, the model identifies patterns that affect pricing. The goal is to enhance decision-making in real estate through data-driven insights, offering value to buyers, sellers, and investors alike.

Index Terms—House Price Prediction, Machine Learning, Real Estate, Regression Models, Random Forest

I. INTRODUCTION

A. Context and Importance of the Real Estate Industry

The real estate sector is a dynamic and vital component of the global economy, influenced by a complex mix of economic, social, and environmental factors. Property transactions typically involve significant financial decisions, making accurate house price prediction essential. Traditional estimation methods often depend on expert opinions or generalized pricing formulas, leading to inconsistencies and subjectivity. These limitations highlight the need for more robust and consistent forecasting mechanisms.

B. Problem Statement

Current approaches to house price estimation primarily rely on manual evaluation or basic statistical techniques. Such reliance on human judgment introduces the risk of error and bias.

Therefore, there is a compelling need for automated, data-driven solutions that can provide more reliable and objective predictions.

C. Role of Machine Learning in Real Estate

With the advent of big data and improvements in machine learning, predictive modeling has become a powerful approach in real estate analytics. Machine learning models can identify patterns from large volumes of historical housing data, enabling more accurate and efficient forecasts. This data-driven insight significantly enhances the decision-making process in real estate transactions.

D. Project Objective

The core objective of this project is to develop a machine learning-based house price prediction system. The system is intended to analyze historical property data and use relevant features to forecast future housing prices accurately. This tool will empower various stakeholders to make informed decisions.

E. Key Features and Influencing Variables

The prediction model incorporates several key attributes that influence property values:

- Location
- Square footage
- Number of rooms (bedrooms and bathrooms)
- Year built
- Proximity to schools, hospitals, and public transport
- Neighborhood quality and crime rates

F. Machine Learning Models Used

The project employs and compares multiple regression algorithms, including:

- Linear Regression
- Random Forest Regressor
- Gradient Boosting Regressor

These models are chosen for their ability to learn complex data relationships and generalize well to new, unseen examples.

G. Data Preprocessing and Engineering

Data preprocessing steps are critical to ensure model performance. These include:

- Handling missing values
- Feature normalization
- Encoding categorical variables (e.g., one-hot encoding)

Feature engineering enhances the dataset by creating meaningful variables that improve model accuracy.

H. Evaluation Metrics

To assess model performance, the following standard metrics are used:

- Mean Absolute Error (MAE)
- Root Mean Squared Error (RMSE)
- Coefficient of Determination (R^2 Score)

I. Technologies and Tools Used

The following technologies support system development and experimentation:

- **Language:** Python
- **Libraries:** Pandas, Scikit-learn, Matplotlib
- **Datasets:** Boston Housing Dataset, Kaggle House Prices Dataset

J. Expected Outcomes

The final product will be a robust, user-friendly house price prediction system capable of:

- Delivering accurate price estimates based on user-input features
- Enhancing transparency and fairness in the housing market
- Assisting buyers, sellers, and investors in making data-informed decisions

II. LITERATURE REVIEW

House price prediction has been a significant area of research in real estate and data science, driven by the need for accurate property valuation. Over the years, various statistical and machine learning techniques have been explored to enhance the prediction accuracy of house prices using diverse datasets.

Early studies primarily used Hedonic Pricing Models, which consider how different property features (e.g., size, location, number of rooms) affect market value. Although effective to some extent, these models often failed to capture complex nonlinear relationships among variables.

With advancements in computational power and data availability, machine learning algorithms have gained popularity in this domain. For instance, Linear Regression remains a baseline model due to its simplicity and interpretability. However, it assumes a linear relationship between inputs and outputs, which limits its performance in complex datasets.

Recent studies have explored more sophisticated models such as Decision Trees, Random Forests, and Gradient Boosting Machines (GBM). These ensemble methods are known

for their high accuracy and robustness against overfitting. Support Vector Machines (SVM) and K-Nearest Neighbors (KNN) have also been employed, especially for smaller or well-preprocessed datasets.

The Kaggle House Prices dataset, a benchmark for this type of prediction, has been used extensively in literature to test and compare various machine learning approaches. Research by Kumar et al. (2020) showed that Gradient Boosting and XGBoost consistently outperformed traditional methods. Additionally, deep learning models, such as Artificial Neural Networks (ANNs), are being investigated for their ability to learn high-level representations from large datasets.

Preprocessing techniques such as feature engineering, normalization, and handling missing data have been emphasized across studies as critical steps for improving model performance.

Overall, the literature suggests that hybrid and ensemble methods offer the most promise for accurate house price prediction, especially when combined with thorough data preprocessing and feature selection strategies. Building upon the insights gathered from previous research, this project aims to apply and enhance existing machine learning methodologies to predict house prices with greater accuracy and reliability. While the literature showcases a variety of models and approaches, there is still room for improvement in terms of data preprocessing techniques, feature selection, model tuning, and evaluation metrics.

This work leverages the strengths of established models like Linear Regression, Random Forest, and Gradient Boosting, while exploring the potential of ensemble methods to improve predictive performance. The implementation focuses not only on accuracy but also on scalability and interpretability, ensuring that the resulting system can be practically applied in real-world real estate scenarios.

Moreover, the project is designed to bridge the gap between theoretical research and practical application. It incorporates a user-friendly interface and a systematic approach to data handling, training, and testing. By using publicly available datasets and open-source tools, the system is designed to be reproducible, transparent, and adaptable.

The subsequent sections detail the system analysis, design, and implementation processes, followed by a discussion of the results obtained. The goal is to contribute to the growing body of work in predictive analytics and provide a solid foundation for future enhancements in the field of property valuation.

III. SYSTEM ANALYSIS

The system analysis for the house price prediction project involves understanding and defining the technical and functional aspects necessary for building a reliable and efficient prediction model. It includes identifying the problem, analyzing user requirements, selecting appropriate tools, and designing a system that delivers accurate results.

A. Problem Definition:

The main problem addressed is the inaccuracy and inefficiency in manually estimating house prices. Real estate

stakeholders often face difficulties due to market volatility and human error. The aim is to develop a data-driven system that predicts house prices based on key features such as location, area, number of rooms, year built, and amenities.

B. Existing System:

Traditional real estate valuation relies on expert judgment, comparative market analysis, or basic statistical models. These methods are prone to subjectivity and fail to consider complex relationships in data. They often lack scalability and adaptability to different regions or market conditions.

C. Proposed System:

The proposed system uses machine learning algorithms to predict house prices more accurately. It takes a dataset of past house sales and learns patterns from it. Models such as Linear Regression, Random Forest, and XGBoost are employed to train and evaluate performance. The system provides users with a predicted price based on the input features.

D. Functional Requirements:

- Data preprocessing and cleaning.
- Feature selection and engineering.
- Model training and evaluation.
- User interface for input and result display.

E. Non-Functional Requirements:

- High prediction accuracy.
- Scalability for large datasets.
- User-friendly interface.
- Low latency in generating results.

F. Feasibility Study:

- Technical Feasibility: Easily implementable with Python, Jupyter Notebook, and machine learning libraries like scikit-learn.
- Economic Feasibility: Cost-effective due to the use of open-source tools.
- Operational Feasibility: Simple to operate and can be used by real estate professionals or individual buyers.

IV. SYSTEM DESIGN

System design is a crucial phase in the development of the House Price Prediction system. It involves planning the structure of the system, selecting appropriate technologies, and defining the flow of data to ensure functionality, accuracy, and usability. The system is designed in a modular fashion, consisting of multiple interrelated components that work together to predict housing prices based on user input and historical data.

A. Architectural Design

The system follows a three-tier architecture:

- Presentation Layer: The front-end interface that allows users to input house features (e.g., location, area, number of bedrooms, bathrooms).
- Application Layer: The backend logic that processes the inputs, invokes the trained machine learning model, and returns the predicted price.
- Data Layer: The storage and retrieval of data, including the historical dataset and model files.

B. Modules of the System

- Data Collection Module: Collects data from sources such as Kaggle or real estate websites.
- Data Preprocessing Module: Cleans data by handling missing values, encoding categorical features, and normalizing numerical values.
- Feature Selection Module: Identifies the most influential features affecting house prices using correlation or importance metrics.
- Model Training Module: Trains various machine learning models like Linear Regression, Decision Tree, or Random Forest.
- Prediction Module: Uses the trained model to predict the price based on new input.
- User Interface Module: Web or desktop-based UI for users to interact with the system.

C. Technologies Used

- Back-end: Python.
- Machine Learning: scikit-learn, pandas, NumPy, matplotlib for visualization.

V. IMPLEMENTATION

The implementation phase involves putting the system design into action by coding and integrating various components to develop a fully functional house price prediction application. The system is implemented using Python due to its simplicity and extensive libraries for data analysis and machine learning.

A. Environment Setup

Tools and Libraries:

- Python 3.x
- Jupyter Notebook / VS Code
- Libraries: pandas, numpy, scikit-learn, matplotlib, seaborn, joblib
- Optional: Flask or Streamlit for building the web interface

B. Data Preprocessing

- Load dataset (e.g., CSV file from Kaggle)
- Handle missing values (e.g., fill with mean or remove rows)
- Encode categorical variables using one-hot encoding or label encoding

- Normalize/scale numerical features using StandardScaler or MinMaxScaler
- Split the dataset into training and testing sets (e.g., 80

C. Model Training

- Implement machine learning models like:
 - Linear Regression
 - Decision Tree
 - Random Forest
- Train models on training data using `model.fit(X_train, y_train)`
- Evaluate models using metrics like Mean Squared Error (MSE), Root Mean Squared Error (RMSE), and R^2 score

D. Model Selection and Saving

- Compare model performances and select the best-performing one
- Save the model using joblib or pickle for reuse without retraining

E. Web Interface (Optional)

- Create a simple web interface using Flask or Streamlit
- Accept user inputs for features like location, size, number of bedrooms
- Load the saved model and display the predicted house price

F. Testing

- Test the system with sample inputs to ensure accuracy and usability
- Perform validation to ensure the model generalizes well

This step-by-step implementation results in an end-to-end system that can predict house prices based on user-defined parameters, offering insights for real estate investors and home buyers.

VI. RESULT AND DISCUSSION

The House Price Prediction system was developed using multiple machine learning algorithms to evaluate their performance and effectiveness in accurately predicting housing prices. The key models applied were Linear Regression, Decision Tree, and Random Forest. After training and testing these models on a cleaned and preprocessed housing dataset, their performances were compared using standard evaluation metrics such as R^2 Score, Mean Absolute Error (MAE), and Root Mean Square Error (RMSE).

Among all models, Random Forest provided the highest prediction accuracy. It achieved an R^2 Score of 0.87, indicating that it could explain 87

A Predicted vs Actual Price graph was plotted for all models, where the Random Forest model showed the closest alignment between predicted and actual prices, affirming its strong generalization capability. The feature importance plot from Random Forest revealed that location, square footage, number of bedrooms, number of bathrooms, and property age were the most influential factors in determining house prices.

The results suggest that ensemble learning models like Random Forest are better suited for this type of regression problem, especially when the data includes a mix of numerical and categorical variables and non-linear patterns. The use of feature engineering and proper data preprocessing played a vital role in improving model performance.

However, despite promising results, there are a few limitations. The model's accuracy depends heavily on the quality and quantity of data. Biases in the dataset can affect prediction accuracy, and regional differences may not be captured effectively if the dataset lacks diversity. Also, external market factors such as inflation, interest rates, and economic conditions were not considered, which could impact real-world pricing.

For future work, the system could be enhanced by integrating real-time data sources, using deep learning models, or deploying the model as a web application for public use. Further tuning and validation using cross-validation methods would also help improve robustness.

Following the detailed analysis of the results and discussion, it becomes essential to contextualize the findings with respect to the project's goals and practical applications. The house price prediction model, based on machine learning techniques, demonstrated promising accuracy and efficiency when tested against real-world housing datasets. These results indicate that such predictive systems can play a significant role in assisting buyers, sellers, and real estate professionals in making informed decisions. By analyzing features such as location, number of bedrooms, square footage, and other key variables, the model effectively captured market trends and provided relatively accurate price estimations.

However, while the model's performance was satisfactory, certain limitations were observed. Variations in housing prices due to economic fluctuations, unstructured data, and missing values in the dataset occasionally affected the prediction accuracy. These challenges point to the importance of enhancing the data quality and incorporating external factors such as market demand, local economic conditions, and seasonal trends in future iterations of the model.

The evaluation metrics, including Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and R-squared values, provided quantifiable evidence of the model's performance. These metrics not only validated the approach but also offered a framework for future comparisons.

Overall, the insights gained from the results support the feasibility of integrating machine learning models into real estate platforms. This can ultimately lead to smarter, data-driven property pricing mechanisms that benefit various stakeholders. The findings pave the way for drawing meaningful conclusions and considering improvements in future work.

VII. CONCLUSION

The objective of this project was to develop an efficient and reliable house price prediction system using machine learning algorithms. Through the application of data preprocessing techniques and model training, the system was able to estimate house prices based on various features such as location, number of rooms, area, and other property-related parameters. The implementation of machine learning models—such as Linear Regression, Decision Tree, and Random Forest—enabled a comparative study of performance metrics to identify the most accurate prediction model.

Among the models tested, the Random Forest Regressor delivered the most accurate results, achieving a high R^2 score and minimal mean squared error (MSE). This demonstrates its robustness in handling complex datasets and its ability to model nonlinear relationships between features. The model also performed well in avoiding overfitting, a common issue in predictive systems, thus ensuring generalization to unseen data.

One of the significant outcomes of the project is the identification of the most influential factors affecting house prices. Features like location, total square footage, number of bedrooms, and age of the property were found to contribute significantly to price variation. This insight not only aids in accurate predictions but also helps real estate stakeholders make informed decisions.

The system has potential applications in real estate platforms, where buyers and sellers can use it to estimate property values more accurately. It can also serve as a decision-support tool for real estate investors and analysts looking to identify undervalued or overpriced properties.

While the project has met its primary objectives, it also opens avenues for further enhancement. Incorporating real-time data sources, expanding the dataset to include multiple cities, and integrating deep learning models could further improve prediction accuracy. Additionally, deploying the system as a user-friendly web or mobile application would make it accessible to a broader audience.

In conclusion, this project successfully demonstrates how machine learning can be effectively applied to predict house prices, offering both technical insight and real-world utility. Future work will focus on enhancing system capabilities and scalability for broader use.

VIII. FUTURE WORK

While the current house price prediction system has demonstrated promising results, there is considerable scope for improvement and expansion. The following points outline potential directions for future development:

Integration of More Features: Future iterations can incorporate additional features such as proximity to amenities (schools, hospitals, transport), crime rates, and neighborhood growth trends, which can significantly influence property values.

Real-time Data Updates: Linking the system to real-time databases, property listing APIs, or government real estate

datasets would enable dynamic predictions based on the most recent market data.

Advanced Algorithms: Implementing deep learning models such as Artificial Neural Networks (ANNs) or Gradient Boosting Machines (e.g., XGBoost, LightGBM) could enhance prediction accuracy, especially in capturing complex nonlinear relationships.

Geospatial Analysis: Using Geographic Information Systems (GIS) data and techniques could provide spatial insights, such as heatmaps of property values, helping users visually interpret market trends.

Cross-city or Country-wide Analysis: Expanding the dataset to cover multiple cities or even countries would make the model more versatile and applicable to a broader range of users.

User Interface and Deployment: Developing a user-friendly web or mobile application would make the system more accessible. This includes features like map integration, interactive dashboards, and personalized prediction reports.

Model Interpretability: Integrating explainable AI (XAI) methods like SHAP or LIME would help users understand how different features impact predictions, increasing trust in the system.

Handling Market Fluctuations: Incorporating economic indicators and seasonality factors can improve resilience against market volatility, making predictions more reliable over time.

Recommendation System: Enhancing the system to not only predict prices but also suggest optimal buying/selling times or investment opportunities based on predicted market trends.

By addressing these areas, the house price prediction system can evolve into a comprehensive real estate decision-support tool with high practical value and adaptability. In addition to the proposed enhancements, future work can focus on incorporating deep learning techniques such as convolutional neural networks (CNNs) or recurrent neural networks (RNNs) to capture complex patterns in spatial and temporal real estate data. Integrating external factors like social media trends, regional development plans, and real-time economic indicators may also help refine prediction models.

Moreover, creating an interactive platform or mobile application that allows users to input custom data (e.g., property features, neighborhood preferences) can provide personalized price estimations and insights. This system could be further enhanced by enabling integration with GIS mapping to offer a visual representation of pricing trends across regions.

Another potential area for exploration is the application of transfer learning, where pre-trained models are adapted to new regions or property types with minimal data. This would be especially useful in areas with limited historical pricing data.

Finally, ethical considerations such as algorithmic bias and data privacy must be addressed in future versions of the system to ensure fairness, transparency, and user trust. Implementing explainable AI (XAI) techniques will help users understand how predictions are made, fostering greater confidence in the model's outputs.

ACKNOWLEDGEMENT

I would like to express my sincere gratitude to all those who supported and guided me throughout the course of this project.

First and foremost, I extend my heartfelt thanks to my project guide, [Insert Guide's Name], for their invaluable insights, encouragement, and continuous support during every phase of the research. Their constructive feedback and expert guidance played a crucial role in the successful completion of this work.

I am also thankful to the faculty members of the Computer Application at KIET Group of Institutions for providing a solid academic foundation and access to the necessary resources and facilities.

Special thanks to my classmates and friends for their encouragement, suggestions, and moral support, which helped me stay motivated throughout the project.

I would also like to acknowledge the open-source community and various online platforms that provided tools, datasets, and reference materials which were essential for the implementation and validation of the system.

Finally, I am deeply grateful to my family for their constant encouragement, patience, and understanding, which enabled me to focus and work diligently on this project.

Without the contributions and support of all these individuals, this project would not have been possible..

REFERENCES

- [1] Kumar, A., Singh, M. (2018). House Price Prediction using Machine Learning Algorithms. *International Journal of Computer Applications*, 179(9), 14-18.
- [2] Li, Y., Li, M. (2017). Housing Price Forecasting Based on Regression Algorithms. *Procedia Computer Science*, 122, 872–876.
- [3] Zhang, D., Chen, L. (2019). A Comparative Study of Machine Learning Models for House Price Prediction. *IEEE Access*, 7, 170353–170362.
- [4] Harrison, D., Rubinfeld, D. L. (1978). Hedonic Prices and the Demand for Clean Air. *Journal of Environmental Economics and Management*, 5(1), 81–102. (Source of the Boston Housing Dataset)
- [5] Scikit-learn Developers. (2024). Scikit-learn: Machine Learning in Python. Retrieved from <https://scikit-learn.org/>
- [6] Kaggle. (2024). House Prices: Advanced Regression Techniques Dataset. Retrieved from <https://www.kaggle.com/c/house-prices-advanced-regression-techniques>
- [7] Ge'ron, A. (2019). *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow* (2nd ed.). O'Reilly Media.
- [8] Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.