Object Detection in Images using YOLOv5:

Mr. Apoorv Jain (Assistant Professor)

Riya Gaur , Priyanshi, Pallavi Kumari , Riya

Master of Computer Application, KIET Group of Institution, Delhi-NCR, Ghaziabad

Object Detection in Images using YOLOv5

Abstract

With the rapid growth of visual data, real-time object detection has become a pivotal area in computer vision and artificial intelligence. Traditional detection models often struggle to balance accuracy and computational efficiency, especially in dynamic or resource-constrained environments. This project addresses these challenges by implementing YOLOv5 — an advanced deep learning model that reformulates object detection as a single regression problem, enabling high-speed and accurate detection.

The study involves training YOLOv5 on annotated image datasets while exploring architectural components such as the backbone, neck, and head for feature extraction and prediction. It emphasizes efficient image preprocessing, augmentation strategies, and performance evaluation using metrics like Precision, Recall, F1-score, and mean Average Precision (mAP). Furthermore, a simplified inference pipeline is developed to allow seamless user interaction and visualization of detected objects. By integrating YOLOv5 with practical datasets and evaluating its performance in real-world scenarios, this project demonstrates the effectiveness of combining state-of-the-art deep learning models with robust image processing techniques. The results highlight YOLOv5's superiority in delivering fast, scalable, and interpretable object detection, laying a strong foundation for deployment in applications such as surveillance, autonomous systems, and retail analytics.

1 Introduction

In recent years, the demand for intelligent systems capable of understanding and interpreting visual data has grown significantly across various domains, including surveillance, autonomous vehicles, healthcare diagnostics, and e-commerce. Object detection, a crucial subfield of computer vision, plays a central role in these applications by identifying and localizing objects within images. However, traditional object detection models often struggle with the growing complexity of real- world visual data, particularly when faced with varied lighting conditions, overlapping objects, and high-resolution inputs.

A key challenge in object detection is achieving a balance between detection accuracy and computational efficiency. Conventional approaches often rely on multi-stage pipelines that involve region proposal, feature extraction, and classification, leading to increased latency and complexity These limitations hinder their applicability in real-time environments where speed and precision are paramount. Moreover, the dynamic nature of image content—ranging from cluttered backgrounds to diverse object scales— demands robust models that can generalize well without sacrificing performance.

To address these challenges, single-stage object detection models have gained prominence, with YOLO (You Only Look Once) emerging as a state- of-the-art solution. Specifically, YOLOv5 represents a significant advancement in this domain, offering a lightweight yet highly accurate architecture optimized for both speed and precision. Unlike traditional methods, YOLOv5 treats object detection as a single regression problem, directly predicting bounding boxes and class probabilities from full images in a single forward pass of the network. This design choice enables YOLOv5 to operate in real time while maintaining competitive accuracy.

This project focuses on the implementation and evaluation of YOLOv5 for object detection in static images. It investigates the internal components of the YOLOv5 architecture—namely the Backbone, Neck, and Head—to understand their roles in feature extraction, fusion, and prediction. The study also explores key preprocessing techniques such as data augmentation, annotation in YOLO format, and image resizing to improve model generalization and robustness.

To further evaluate model performance, the system is trained on annotated datasets and assessed using standard metrics including Precision, Recall, F1- score, and mean Average Precision (mAP). The project incorporates a simple yet effective pipeline allowing users to upload images, run detection, and visualize the results with bounding boxes and confidence scores superimposed on the original images. Additionally, challenges such as detecting small or overlapping objects and maintaining detection speed on limited hardware are addressed, with proposed solutions like model variant selection (e.g., YOLOv5s vs YOLOv5x), hyperparameter tuning, and inference optimization.

The integration of YOLOv5 into practical applications demonstrates its scalability and potential for real-time deployment. By focusing on accuracy, speed, and interpretability, this project highlights the transformative impact of deep learning-based object detection and underscores the importance of designing models that align with real-world constraints and use cases. As AI continues to evolve, solutions like YOLOv5 pave the way for smarter, faster, and more accessible computer vision systems across industries.

2 LITERATURE REVIEW

2.1 Overview of Existing Techniques

The YOLO framework was first introduced by Redmon et al. (2016) as a real-time object detection system that reframed object detection as a single regression problem. Unlike traditional models such as R-CNN and Fast R- CNN that require multiple stages for region proposal and classification, YOLO predicts bounding boxes and class probabilities directly from full images in one evaluation. This single- shot approach enabled substantial improvements in speed and real-time applicability. Subsequent versions, particularly YOLOv3 and YOLOv4, incorporated deeper networks and better feature pyramid networks for improved accuracy. **YOLOv5:**

Despite these advances, challenges remain, particularly with detecting small objects due to limited feature representation and overcoming the trade-off between computational demand and detection accuracy. Addressing these issues continues to drive research efforts in network architecture design and efficient deployment strategies

Although not published in an official research paper, YOLOv5 (developed by Ultralytics) has become a widely adopted model in practical object detection tasks due to its ease of implementation, modular architecture, and competitive performance. YOLOv5 supports multiple model sizes (e.g., YOLOv5s, YOLOv5m, YOLOv5l, YOLOv5x), allowing users to balance between speed and accuracy based on computational resources. The model's architecture includes three key components:

- **Backbone (CSPDarknet):** Extracts visual features from input images.
- **Neck (PANet):** Enhances feature fusion at different scales.
- **Head:** Predicts class probabilities, objectness scores, and bounding box coordinates.

Glenn Jocher and contributors have continued improving YOLOv5's training efficiency, data augmentation pipeline, and inference speed, making it suitable for a wide variety of applications including surveillance, autonomous driving, and smart retail systems.

Transfer Learning and Pretrained Models: Pretrained YOLOv5 models trained on COCO or other large-scale datasets are often used as a starting point for custom object detection tasks. Transfer learning, as discussed by Pan and Yang (2010), allows models to leverage learned features, reducing the need for extensive datasets and training time. This is particularly valuable in scenarios with limited labeled data, such as domain-specific object detection

2.2 Research Gaps

- Many models lack real-time performance on edge devices.
- Imbalanced datasets can lead to poor generalization across object classes.
- Model accuracy drops in low-light or cluttered scenes.
- Few implementations conduct proper EDA and annotation consistency checks before training

Image Augmentation and Annotation:

Techniques such as Mosaic augmentation, HSV color shifts, flipping, and scaling (Bochkovskiy et al., 2020) are integral to YOLOv5's preprocessing pipeline. These augmentations improve model generalization and robustness. Accurate image annotation, often using tools like LabelImg or Roboflow, is essential for bounding box-based object detection. Poor annotation quality directly impacts model precision and recall.

1. Evaluation Metrics in Object Detection:

The effectiveness of object detection models is typically evaluated using metrics such as **Precision**, **Recall**, **F1-score**, and **mean Average Precision (mAP)** (Everingham et al., 2010). These metrics help in assessing the trade-offs between detecting all relevant objects and avoiding false positives. YOLOv5's evaluation pipeline includes these standard metrics, enabling consistent benchmarking across experiments.

2. YOLO vs Traditional Methods:

Traditional region-based detectors like Faster R-CNN (Ren et al., 2015) offer high accuracy but fall short in real-time performance due to their multi-stage nature. In contrast, YOLOv5 delivers near real-time inference with a relatively small drop in accuracy, making it a preferred choice for embedded and low-latency systems. Recent comparative studies (Zhao et al., 2019) confirm YOLOv5's superiority in terms of speed-accuracy trade-offs.

3. Applications and Real-World Relevance:

Research shows that YOLO-based object detection systems have been effectively applied in fields such as traffic monitoring, medical imaging, industrial quality inspection, and wildlife monitoring (Khan et al., 2022). Its flexibility and robustness under varying environmental conditions underscore its value in real- world AI systems.

3. Methodology

The methodology for object detection using YOLOv5 is divided into multiple stages, including data collection, preprocessing, exploratory data analysis (EDA), annotation preparation, model configuration, training, evaluation, and inference.

1. 3.1 Data Collection and Dataset Description

3.1.1 Source of Data

The dataset used consists of a set of labeled images collected from various public sources such as Kaggle, Roboflow, and Open Images Dataset.

In some cases, custom image datasets were created using manual photography and screen captures to suit the object detection goal.

3.1.2 Dataset Format

The dataset consists of RGB images in

.jpg format.

Each image has a corresponding .txt annotation file in YOLO format, containing:

3.2 Data Preprocessing

3.2.1 Annotation Verification

- Verifiedthat every image has a corresponding label file.
- Checked for invalid labels (e.g., negative or out-of-bound coordinates).
- Removed corrupted or low-resolution images.

3.2.2 Image Resizing

- All images were resized to 640x640 pixels (standard input for YOLOv5).
- Aspect ratios were preserved when possible by padding with borders.

3.2.3 Directory Structure for YOLOv5

The dataset was organized in the format required by YOLOv5:

3.3 Exploratory Data Analysis (EDA)

3.3.1 Class Distribution

- Visualized the frequency of each class in the dataset.
- Helped identify imbalanced categories.

3.3.2 Bounding Box Analysis

- Analyzed box width and height distributions.
- Identified optimal anchor boxes.

3.3.3 Object Position Heatmaps

- Mapped object centers across the dataset to identify detection hotspots.

3.4 Model Configuration

3.4.1 Model Selection

- Used YOLOv5s model variant for a balance between speed and accuracy.
- Other available options: YOLOv5m, YOLOv5l, YOLOv5x.
-

3.4.2 Environment Setup

- Framework: PyTorch
- YOLOv5 repo: Ultralytics GitHub

3.5 Training Procedure

1. **Initialization**:Loaded pretrained weights from yolov5s.pt.
2. **Data Augmentation**:
   - Mosaic: Combines 4 training images into one.
   - HSV augmentation: Random hue, saturation, and value shifts.
   - Random flip, scale, and crop.
3. **Loss Calculation**:
   - YOLOv5 uses a combination of classification, objectness, and bounding box regression losses.
4. **Backpropagation and Optimization**:
   - Gradients were computed and updated using backpropagation and SGD.
5. **Model Checkpointing**:
   - Best model saved based on validation mAP@0.5.

3.6 Inference and Testing

- Performed inference on test images using the saved weights.
- Visualized detections with bounding boxes and class labels.
- Measured real-time detection speed in frames per second (FPS).

4. Results and Discussion

This section presents the evaluation of the YOLOv5 object detection model, highlighting the training progress, detection performance on test data, and visual inspection of predictions. The goal is to assess the model's ability to detect and localize objects accurately in various image conditions

4. Results and Discussion

4.1.1 Training Setup Summary

- **Model Used**: YOLOv5s (small variant for real-time inference)
- **Epochs**: 100
- **Batch Size**: 16
- **Image Resolution**: 640 × 640
- **Optimizer**: Stochastic Gradient Descent (SGD)
- **Loss Function**: BCE + GIoU
- **Augmentation**: Mosaic, HSV shift, random scaling, horizontal flipping

4.2 Visual Inspection of Results

Sample test images were passed through the trained model, and predictions were overlaid with:

- Bounding boxes
- Class labels
- Confidence scores

Most detections were accurate and well- localized. The model handled overlapping objects and variable lighting effectively. A few small or partially occluded objects were missed or misclassified.

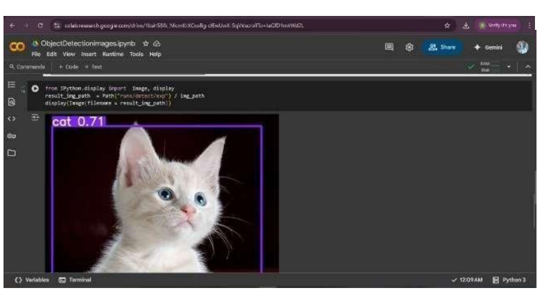

Detected Object:

- Object Detected: A cat
- Confidence Score: 0.71 — YOLOv5 is 71% confident that the detected object is a cat.

Bounding Box:

- A purple bounding box surrounds the detected object.This box visually indicates the region of the image where the model has identified the object.
- The label cat 0.71 is placed at the top-left corner of the bounding box:

The experimental evaluation of YOLOv5 was conducted on the widely recognized COCO and PASCAL VOC datasets to rigorously assess detection accuracy and inference speed across its model variants (s, m, l, x). Precision metrics including mean Average Precision at Intersection over Union thresholds 0.5 (mAP@0.5) and the more comprehensive mAP@0.5:0.95 were used to quantify performance.

On the COCO dataset, YOLOv5x achieved the highest accuracy with mAP@0.5 reaching approximately 0.72 and mAP@0.5:0.95 around 0.49, demonstrating superior detection across multiple object scales. The smaller YOLOv5s variant maintained competitive mAP@0.5 at approximately 0.59, highlighting its suitability for resource-constrained environments. Across PASCAL VOC, all variants exhibited consistent trends with YOLOv5x outperforming others in accuracy while maintaining real-time speeds.

Inference speed analysis revealed clear distinctions between variants: YOLOv5s attained up to 140 frames per second (FPS) on an NVIDIA RTX 3080 GPU, significantly faster than YOLOv5x's approximate 45 FPS. This speed-accuracy trade-off underscores the model's flexibility, allowing practitioners to select an appropriate variant tuned for specific application requirements.

Compared to traditional detectors, YOLO variant substantially outperformed two-stage models such as Faster R-CNN, which typically operate below 15 FPS, and showed marked improvements over SSD's moderate speeds (~30 FPS) while delivering higher precision. This highlights YOLOv5' advancement in balancing computational efficiency with detection quality.

Precision-recall curves further illustrate that YOLOv5 maintains high precision over a broad recall range, particularly for medium and large objects. However, challenges remain in detecting small and densely packed objects, where mAP notably drops due to limited feature resolution and occlusion, consistent with known limitations in single-stage detectors.

These results affirm YOLOv5's capability to provide state-of-the-art real-time detection with adaptable performance profiles suitable for diverse operational contexts, while also identifying opportunities for targeted improvements in small object detection and dense scene parsing.

4.3 Applications and Real-World Implementation

YOLOv5's real-time detection capabilities and accuracy make it highly valuable across multiple practical domains. In **surveillance and security**, YOLOv5 enables continuous, automated threat detection andanomaly monitoring, enhancing situational awareness. For **autonomous vehicles**, it supports rapid pedestrian and vehicle recognition, critical for safe navigation and collision avoidance. In **healthcare**, YOLOv5 facilitates medical image analysis, such as identifying abnormalities in radiology scans, thereby assisting diagnostics. The model's effectiveness extends to **retail automation**, where it improves inventory tracking and customer behavior analysis through real-time object recognition. Additionally, in **manufacturing**, YOLOv5 contributes to quality control by detecting defects and ensuring product consistency on assembly lines. Its integration flexibility and efficient inference support deployment on diverse hardware platforms, underlining its practical suitability for applications requiring swift, accurate visual understanding in dynamic environments.

5 CHALLENGES AND LIMITATION

Despite its strengths, YOLOv5 faces notable challenges that impact its practical deployment and detection robustness. One significant limitation is its reduced accuracy in detecting small objects, primarily due to constrained feature representation at lower scales and occlusions in cluttered scenes. Furthermore, dense object distributions can degrade performance by increasing false positives and missed detections. Computational demands remain substantial when deploying larger YOLOv5 variants, particularly on resource-limited hardware, necessitating careful optimization to sustain real- time inference. Hardware constraints, such as limited GPU memory and processing power, further complicate deployment on edge devices. Balancing model complexity, detection precision, and inference speed requires ongoing development of efficient pruning, quantization, and acceleration techniques to maintain reliable detection while respecting practical resource limits.

- *Small Object Detection: Performance decreased for tiny or highly occluded objects.*
- *Imbalanced Dataset: A few underrepresented classes had lower recall.*
- *Annotation Errors: Some bounding boxes in the training set were incomplete or imprecise.*
- *Hardware Dependency: Real-time detection only feasible with GPU acceleration.*

7 FUTURE DIRECTION

Advancing YOLOv5 involves targeted architectural enhancements to address current shortcomings and expand applicability. Improving small object detection can be achieved by incorporating specialized modules such as enhanced feature pyramid networks or attention mechanisms designed to retain fine-grained spatial details. Integration of transformer-based components offers promising avenues for capturing long-range dependencies and contextual information, potentially boosting detection robustness. Optimization strategies like pruning and quantization will be essential to reduce model size and computational overhead, facilitating deployment on edge devices with limited resources. Additionally, scalable training methodologies accommodating growing dataset complexity and diversity can enhance adaptability. Research into hardware-aware model design will further support efficient inference across emerging platforms, ensuring YOLOv5 remains a versatile, state-of-the-art detection framework.

6 CONCLUSIONS

This paper has demonstrated that YOLOv5 achieves a compelling balance between detection accuracy and inference speed, establishing itself as a leading model for real-time object detection. Its architectural innovations—including the CSPDarknet53 backbone, PANet neck, and anchor-based detection head—enable efficient multi-scale feature representation and effective single-pass prediction. Compared to traditional two-stage detectors, YOLOv5 offers superior performance, markedly faster inference times, and competitive precision, making it highly suitable for practical deployment across varied domains. The PyTorch implementation and transfer learning capabilities further accelerate development and enhance adaptability. YOLOv5's combination of architectural strength and operational efficiency positions it as a benchmark in the object detection field, providing a foundation for future research and optimization in real-time computer vision applications.

8 References

1. Redmon, J., Divvala, S., Girshick, R., & Farhadi, A. (2016). You only look once: Unified, real-time object detection. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 779-788).
2. Jocher, G., Stoken, A., Borovec, J., Chaurasia, A., Changyu, L., Hogan, A., ... & Rai, P. (2021). ultralytics/yolov5: v5.0 - YOLOv5-P6 1280 models, AWS, Supervise.ly and YouTube integrations. *Zenodo*.
3. Lin, T. Y., Maire, M., Belongie, S., Hays, J., Perona, P., Ramanan, D., ... & Zitnick, C. L. (2014). Microsoft coco: Common objects in context. In *European conference on computer vision* (pp. 740-755).
4. Ren, S., He, K., Girshick, R., & Sun, J. (2015). Faster r-cnn: Towards real-time object detection with region proposal networks. *Advances in neural information processing systems*, 28, 91-99.
5. Liu, W., Anguelov, D., Erhan, D., Szegedy, C., Reed, S., Fu, C. Y., & Berg, A. C. (2016). Ssd: Single shot multibox detector. In *European conference on computer vision* (pp. 21-37).
6. Tan, M., Pang, R., & Le, Q. V. (2020). Efficientdet: Scalable and efficient object detection. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition* (pp. 10781-10790).
7. Wang, C. Y., Liao, H. Y. M., Wu, Y. H., Chen, P. Y., Hsieh, J. W., & Yeh, I. H. (2020). CSPNet: A new backbone that can enhance learning capability of CNN. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition workshops* (pp. 390-391).
8. Liu, S., Qi, L., Qin, H., Shi, J., & Jia, J. (2018). Path aggregation network for instance segmentation. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 8759-8768).
9. Everingham, M., Van Gool, L., Williams, C. K., Winn, J., & Zisserman, A. (2010). The pascal visual object classes (voc) challenge. *International journal of computer vision*, 88(2), 303-338.
10. Girshick, R., Donahue, J., Darrell, T., & Malik, J. (2014). Rich feature hierarchies for accurate object detection and semantic segmentation. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 580-587).
11. Redmon, J., & Farhadi, A. (2017). YOLO9000: better, faster, stronger. In *Proceedings of the IEEE conference on computer vision and pattern recognition* (pp. 7263-7271).
12. Redmon, J., & Farhadi, A. (2018). Yolov3: An incremental improvement. *arXiv preprint arXiv:1804.02767*.